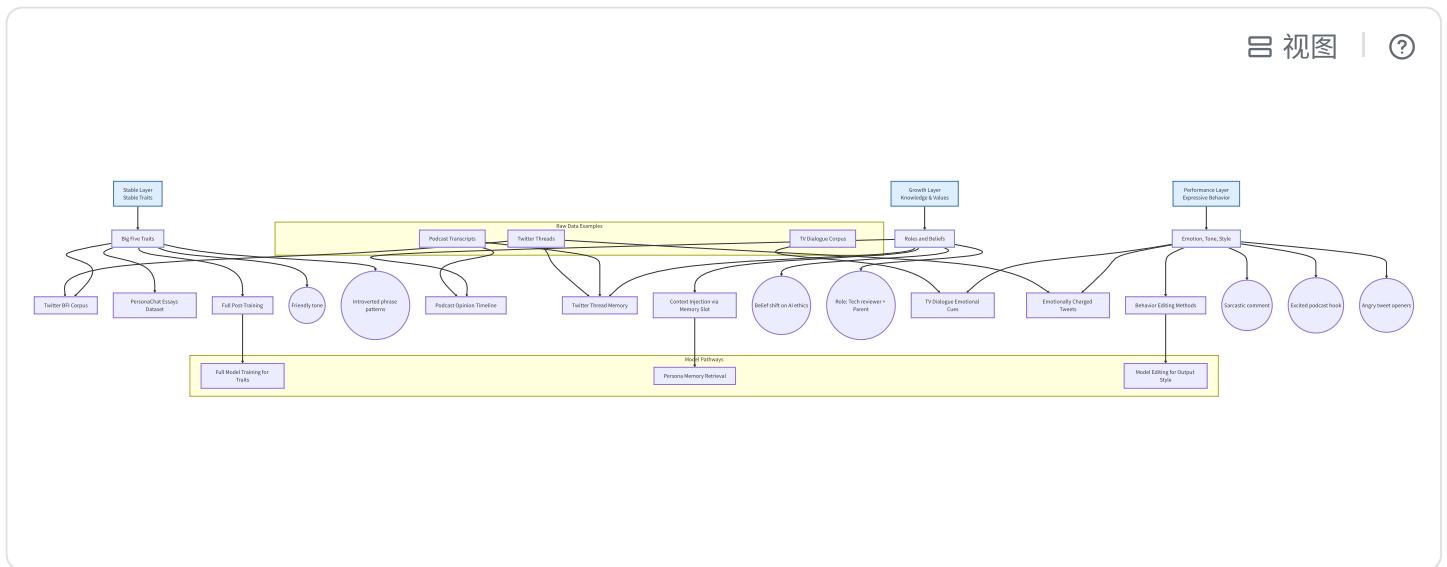


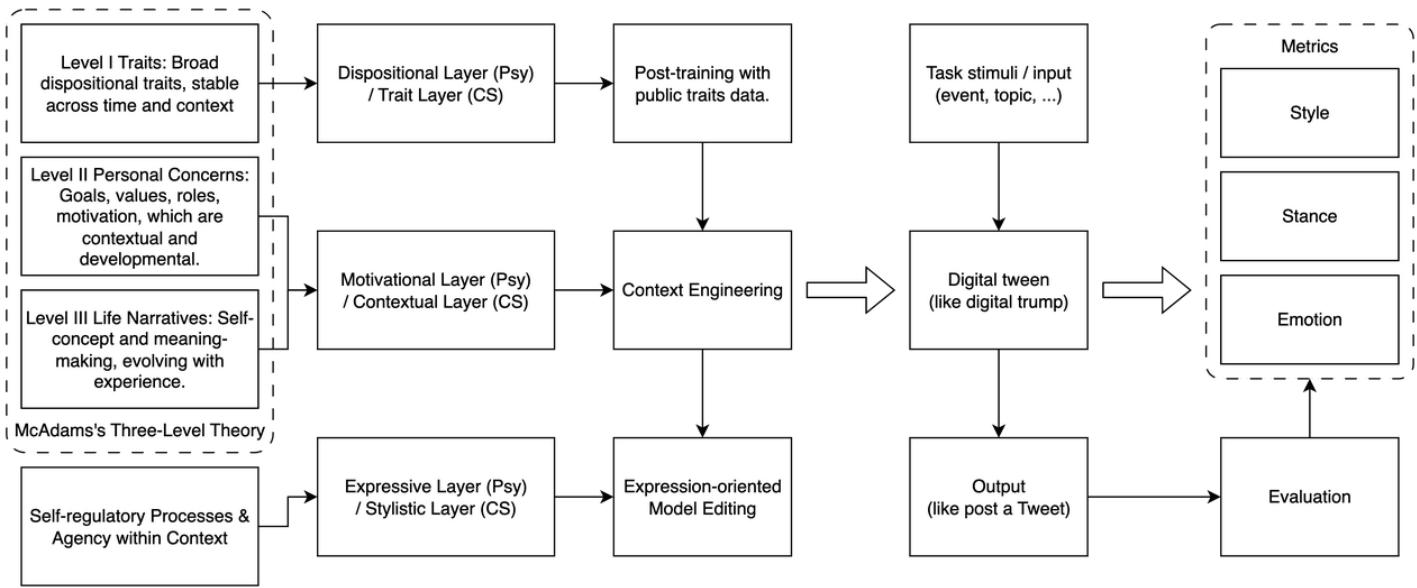
A Three-Tier Personality Framework for Behaviorally Consistent LLMs



Domains and major life events Personality change

Occupational domain	
Graduation from school	C+, ES+, E-
Work life entry	C+, ES+
Job change	E and O (direction depends on role demand)
Promotion at work	C+, ES+, O+
Unemployment	C-, O-, A-
Retirement	C-
Social domain	
Starting a romantic relationship	ES+, E+, A-
Marriage	ES+, E-, O-, A-
Separation/divorce	O+, A+
Birth of a child	E-, C-
Health domain	ES-, E-, O-, C-

The impact of major life events on the Big Five personality traits. ES, emotional stability; E, extraversion; O, openness to experience; A, agreeableness; C, conscientiousness; +, tends to increase due to the event; -, tends to decrease due to the event.



An Example: Digital Trump

[Building Digital Trump: A Layered Personality Modeling and Behavior Simulation Pipeline](#)

1. Background and Motivation

Large Language Models (LLMs) are increasingly being deployed in interactive and multi-role scenarios, where coherence, personality, and social alignment are key to user trust and engagement. Despite recent progress in persona-conditioned generation, current techniques—such as prompt injection, fine-tuning, or memory augmentation—often fall short in maintaining consistency across contexts and over time. These approaches tend to treat “personality” as a surface-level token control or style shift, lacking a principled and structured framework that reflects how personality governs behavior in humans.

In contrast, psychological science offers well-established models that distinguish between relatively stable personality traits (e.g., the Big Five), contextual adaptations (e.g., values, goals, and motivations), and moment-to-moment behavioral expressions (e.g., emotions, tone, stance). Frameworks such as McAdams' three-layer personality model, the Five Factor Theory (FFT), and the TESSERA framework describe personality as a layered and dynamic system. However, these insights have yet to be systematically integrated into LLM design.

This project aims to fill that gap by proposing a three-tier personality modeling framework for LLMs, grounded in psychological theory:

- **Stable Layer:** enduring dispositions such as personality traits and identity;
- **Growth Layer:** experience-based adaptations like goals, motivations, and role-specific knowledge;
- **Performance Layer:** expressive behaviors including tone, stance, affect, and verbal style, modulated by situational context.

2. Research Objectives and Contributions

This research aims to establish a structured and psychologically grounded framework for modeling, injecting, and evaluating personality in LLMs. The key contributions include:

1. **A Three-Tier Personality Framework for LLMs:** A Stable–Growth–Performance architecture supporting enduring consistency and flexible expression.
 2. **Mapping Personality Layers to Model Components:**
 - Stable: trait embeddings/persona profiles;
 - Growth: dynamic memory modules/latent trackers;
 - Performance: prompt-time modulation/vector steering.
 3. **Training and Conditioning Strategies:** SFT, contrastive loss, memory injection, and LoRA steering to inject and control layered personality.
 4. **Evaluation Metrics:**
 - Consistency across contexts;
 - Expressivity in behavior;
 - Adaptability in role-switching.
 5. **Prototype Agent:** A demo LLM agent exhibiting layered personality expression, adaptable over long-term and cross-scenario interactions.
-

3. Methodology and Technical Roadmap

3.1 Layer Specification

- **Stable Layer:** Big Five traits, identity archetypes; fixed vectors or long-term memory entries.
- **Growth Layer:** Acquired values, goals, roles; adaptive modules updated via feedback or exposure.
- **Performance Layer:** Affective states, stance, style; controlled via prompts or dynamic decoding.

3.2 Data Strategy

- Collect multi-context persona-behavior triplets from:
 - Twitter, podcasts, interviews;
 - Persona-Chat, Character-DIAL, RealToxicityPrompts;
 - Public biographies and annotated corpora.
- Use GPT-4/5 for tagging traits, stances, and style labels.

3.3 Model Training

- Base Models: gpt-oss-20b, Qwen3-8b, LLaMA3.1-8b-instruct.
- Techniques:
 - Persona-aligned SFT;
 - Memory-based growth conditioning;
 - Contrastive persona-behavior supervision;
 - Steering module for performance layer.

3.4 Evaluation

- **Consistency:** Persona retention score, trait drift tracking.
- **Expressivity:** Human-style judgment, classification accuracy.
- **Adaptability:** Role-change plausibility, response plausibility.

3.5 Implementation Resources

- **Compute:** A100 GPU x 4–8; ~1000–2000 GPU hours.
 - **LLM API Use:** GPT-4/5 ~15M tokens for preprocessing.
 - **Personnel:** 1 NLP researcher (full-time), 1 data engineer, 1 psychology advisor.
-

4. Related Work

Recent studies on LLM personality modeling include:

- **Prompt-based conditioning:** e.g., Wang et al. (2023), Liu et al. (2023);
- **Emergent persona tracking:** e.g., Zhao et al. (2025);
- **Behavioral injection:** e.g., Li et al. (2025);
- **Memory/vector personalization:** PersonalGPT (Cao et al., 2023), PersonaAdapt (Chen et al., 2024);
- **Psychological alignment:** Caron & Srivastava (2022), Mills & Emmons (2023).

None of these methods integrates multi-layer personality theory or behaviorally grounded modeling across contexts and time.

5. Theoretical Foundations

Our framework is rooted in:

- **McAdams' Three-Level Personality Theory:**

- **Five-Factor Theory (FFT);**
- **Whole Trait Theory (WTT);**
- **TESSERA Framework.**

These models support the stratification of traits, experiences, and behavioral expression, directly informing our LLM component design.

6. Significance and Future Directions

This work offers:

- A principled foundation for interpretable and controllable personality modeling in LLMs;
- A methodology for layered persona injection and adaptation;
- A bridge between psychological theory and machine learning practice.

Future extensions include personality development modeling, real-time emotional feedback integration, and multi-agent personality interaction simulations.